

# Parallelization of a fast Poisson solver

Arne Morten Kvarving

Department of Mathematical Sciences  
Norwegian University of Science and Technology

March 4, 2010


# Problem overview

- We consider the Poisson problem

$$\begin{aligned} -\nabla^2 u &= f && \text{in } \Omega \\ u &= 0 && \text{on } \partial\Omega \end{aligned}$$

in a square domain  $\Omega = (0, 1) \times (0, 1)$ .

- We discretize using finite differences, more specifically using the five-point stencil with  $n + 1$  points in each spatial direction,



## Problem overview contd.

- This gives us the approximation to our problem as a linear system of equations

$$\mathbf{A}\mathbf{u} = \mathbf{g}.$$

of dimension  $(n-1)^2 \times (n-1)^2$ .

- We solve this problem using diagonalization techniques as

1. Transform

$$\tilde{\mathbf{G}} = \mathbf{Q}^T \mathbf{G} \mathbf{Q}$$

2. Scale

$$\tilde{u}_{i,j} = \frac{\tilde{g}_{ij}}{\lambda_i + \lambda_j}, \quad 1 \leq i, j \leq n-1$$

3. Transform back


$$\mathbf{U} = \mathbf{Q} \tilde{\mathbf{U}} \mathbf{Q}^T.$$

## Problem overview contd.

- Due to the eigenvectors of this specific operator, we can perform multiplications with  $\mathbf{Q}$  and  $\mathbf{Q}^T$  using fast sine transforms; see the lecture notes.
- This was a brief overview of the maths behind the serial code you have been given. We now move on to the real problem at hand; parallelization of this code.


# Parallelization - problem partitioning

- We have to decide how we want to divide our problem between the processes.
- We basically have two choices; either a block strategy like



# Parallelization - problem partitioning

- Alternatively we can use a column/row division like



# Parallelization - problem partitioning

- The *fst* operate along whole columns of the matrix.
- The block approach would require parallelization within the *fst* method. This is something we need to avoid.
- Note that since we have  $n - 1$  columns to divide among  $p$  processes, where  $n$  is a power of two, we cannot assign equal amounts of data to each process in general.
- Important hint: Store the designated sizes in an array!

# Parallelization - transposing the matrix

- Since data is stored in several processes, transposing the matrix is nontrivial. This is where most of the effort needs to be invested.
- Since we don't have an equal amount of data in each process, we have to use `MPI_Alltoallv`.
- `MPI_Alltoallv` can only work on a vector. We hence need to pack the matrix into a suitable vector before calling the function.
- We now consider two simple examples in an attempt to give you the idea of what you have to do.



## Parallelization - 3x3 using 2 processes

- Consider the case of a 3x3 matrix distributed on 2 processes.

$$\begin{array}{c|cc} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \end{array}$$

The size array is  $S = \begin{bmatrix} 1 & 2 \end{bmatrix}$ .

- Transposed we have

$$\begin{array}{c|cc} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{array}$$

## Parallelization - 3x3 using 2 processes

- Scount arrays:

proc 0:

$$\begin{vmatrix} 1 & 2 \end{vmatrix} = \begin{vmatrix} S(1)S(1) & S(1)S(2) \end{vmatrix}$$

proc 1:

$$\begin{vmatrix} 2 & 4 \end{vmatrix} = \begin{vmatrix} S(2)S(1) & S(2)S(2) \end{vmatrix}$$

- Sdispl arrays:

proc 0:

$$\begin{vmatrix} 0 & 1 \end{vmatrix} = \begin{vmatrix} 0 & S(1)S(1) \end{vmatrix}$$

proc 1:

$$\begin{vmatrix} 0 & 2 \end{vmatrix} = \begin{vmatrix} 0 & S(2)S(1) \end{vmatrix}$$

- Rcount and Rdispl is the same arrays.

## Parallelization - 3x3 using 2 processes

- Send buffers:

proc 0	proc 1
1	4
2	7
3	5
	6
	8
	9
- We then perform the call to `MPI_Alltoallv`.

- Receive buffers:

proc 0	proc 1
1	2
4	3
7	5
	6
	8
	9

## Parallelization - 7x7 using 3 processes

- Consider the case of a 7x7 matrix distributed on 3 processes.

1	8	15	22	29	36	43
2	9	16	23	30	37	44
3	10	17	24	31	38	45
4	11	18	25	32	39	46
5	12	19	26	33	40	47
6	13	20	27	34	41	48
7	14	21	28	35	42	49

The size array is  $S = \begin{bmatrix} 2 & 2 & 3 \end{bmatrix}$ .

- Transposed we have

1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31	32	33	34	35
36	37	38	39	40	41	42
43	44	45	46	47	48	49

## Parallelization - 7x7 using 3 processes

- Scount arrays:

proc 0:

$$\begin{vmatrix} 4 & 4 & 6 \end{vmatrix} = \begin{vmatrix} S(1)S(1) & S(1)S(2) & S(1)S(3) \end{vmatrix}$$

proc 1:

$$\begin{vmatrix} 4 & 4 & 6 \end{vmatrix} = \begin{vmatrix} S(2)S(1) & S(2)S(2) & S(2)S(3) \end{vmatrix}$$

proc 2:

$$\begin{vmatrix} 6 & 6 & 9 \end{vmatrix} = \begin{vmatrix} S(3)S(1) & S(3)S(2) & S(3)S(3) \end{vmatrix}$$

## Parallelization - 7x7 using 3 processes

- Sdispl arrays:

proc 0:

$$\begin{vmatrix} 0 & 4 & 8 \end{vmatrix} = \begin{vmatrix} 0 & S(1)S(1) & S(1)S(1) + S(1)S(2) \end{vmatrix}$$

proc 1:

$$\begin{vmatrix} 0 & 4 & 8 \end{vmatrix} = \begin{vmatrix} 0 & S(2)S(1) & S(2)S(1) + S(2)S(2) \end{vmatrix}$$

proc 2:

$$\begin{vmatrix} 0 & 6 & 12 \end{vmatrix} = \begin{vmatrix} 0 & S(3)S(1) & S(3)S(1) + S(3)S(2) \end{vmatrix}$$

- Again Rcount and Rdispl is the same arrays.

# Parallelization - 7x7 using 3 processes

- Send buffers:

proc 0	proc 1	proc 2
1	15	29
2	16	30
8	22	36
9	23	37
3	17	43
4	18	44
10	24	31
11	25	32
5	19	38
6	20	39
7	21	45
12	26	46
13	27	33
14	28	34
		35
		40
		41
		42
		47
		48
		49

# Parallelization - 7x7 using 3 processes

- Receive buffers:

proc 0	proc 1	proc 2
1	3	5
2	4	6
8	10	7
9	11	12
15	17	13
16	18	14
22	24	19
23	25	20
29	31	21
30	32	26
36	38	27
37	39	28
43	45	33
44	46	34
		35
		40
		42
		47
		48
		49